

data-wrangling

February 15, 2022

1 Data Wrangling

Estimated time needed: **30** minutes

1.1 Objectives

After completing this lab you will be able to:

- Handle missing values
- Correct data format
- Standardize and normalize data

Table of Contents

Identify and handle missing values

Identify missing values

Deal with missing values

Correct data format

Data standardization

Data normalization (centering/scaling)

Binning

Indicator variable

What is the purpose of data wrangling?

Data wrangling is the process of converting data from the initial format to a format that may be better for analysis.

What is the fuel consumption (L/100k) rate for the diesel car?

Import data

You can find the “Automobile Dataset” from the following link: <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>. We will be using this dataset throughout this course.

Import pandas

```
[ ]: #install specific version of libraries used in lab
    #! mamba install pandas==1.3.3
    #! mamba install numpy=1.21.2
```

```
[1]: import pandas as pd
    import matplotlib.pyplot as plt
```

Reading the dataset from the URL and adding the related headers

First, we assign the URL of the dataset to “filename”.

This dataset was hosted on IBM Cloud object. Click [HERE](#) for free storage.

```
[2]: filename = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
    ↳IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/Data%20files/auto.csv"
```

Then, we create a Python list headers containing name of headers.

```
[3]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
    ↳ "num-of-doors", "body-style",
    ↳ "drive-wheels", "engine-location", "wheel-base",
    ↳ "length", "width", "height", "curb-weight", "engine-type",
    ↳ "num-of-cylinders",
    ↳ "engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
    ↳ "peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter “names” equal to the Python list “headers”.

```
[4]: df = pd.read_csv(filename, names = headers)
```

Use the method `head()` to display the first five rows of the dataframe.

```
[5]: # To see what the data set looks like, we'll use the head() method.
    df.head()
```

```
[5]:  symboling  normalized-losses      make fuel-type aspiration num-of-doors  \
0         3             ?  alfa-romero    gas      std         two
1         3             ?  alfa-romero    gas      std         two
2         1             ?  alfa-romero    gas      std         two
3         2          164      audi      gas      std         four
4         2          164      audi      gas      std         four

    body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front      88.6  ...        130
1  convertible         rwd         front      88.6  ...        130
2   hatchback         rwd         front      94.5  ...        152
3        sedan         fwd         front      99.8  ...        109
4        sedan         4wd         front      99.4  ...        136
```

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

df.head() As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

How to work with missing data?

Steps for working with missing data:

Identify missing data

Deal with missing data

Correct data format

Identify and handle missing values

Identify missing values

Convert “?” to NaN

In the car dataset, missing data comes with the question mark “?”. We replace “?” with NaN (Not a Number), Python’s default missing value marker for reasons of computational speed and convenience. Here we use the function:

to replace A by B.

```
[8]: import numpy as np

# replace "?" to NaN
df.replace("?", np.nan, inplace = True)
df.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	3	NaN	alfa-romero	gas	std	two	
1	3	NaN	alfa-romero	gas	std	two	

2	1	NaN	alfa-romero	gas	std	two
3	2	164	audi	gas	std	four
4	2	164	audi	gas	std	four

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

Evaluating for Missing Data

The missing values are converted by default. We use the following functions to identify these missing values. There are two methods to detect missing data:

```
.isnull()
.notnull()
```

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
[9]: missing_data = df.isnull()
missing_data.head(5)
```

```
[9]:   symboling  normalized-losses  make  fuel-type  aspiration  num-of-doors  \
0      False                True  False    False    False        False
1      False                True  False    False    False        False
2      False                True  False    False    False        False
3      False                False  False    False    False        False
4      False                False  False    False    False        False
```

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	False	False	False	False	...	False	
1	False	False	False	False	...	False	
2	False	False	False	False	...	False	
3	False	False	False	False	...	False	
4	False	False	False	False	...	False	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	city-mpg	highway-mpg	price
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

[5 rows x 26 columns]

“True” means the value is a missing value while “False” means the value is not a missing value.

Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, “True” represents a missing value and “False” means the value is present in the dataset. In the body of the for loop the method “.value_counts()” counts the number of “True” values.

```
[10]: for column in missing_data.columns.values.tolist():
      print(column)
      print (missing_data[column].value_counts())
      print("")
```

symboling

False 205

Name: symboling, dtype: int64

normalized-losses

False 164

True 41

Name: normalized-losses, dtype: int64

make

False 205

Name: make, dtype: int64

```
fuel-type
False      205
Name: fuel-type, dtype: int64

aspiration
False      205
Name: aspiration, dtype: int64

num-of-doors
False      203
True        2
Name: num-of-doors, dtype: int64

body-style
False      205
Name: body-style, dtype: int64

drive-wheels
False      205
Name: drive-wheels, dtype: int64

engine-location
False      205
Name: engine-location, dtype: int64

wheel-base
False      205
Name: wheel-base, dtype: int64

length
False      205
Name: length, dtype: int64

width
False      205
Name: width, dtype: int64

height
False      205
Name: height, dtype: int64

curb-weight
False      205
Name: curb-weight, dtype: int64

engine-type
False      205
```

Name: engine-type, dtype: int64

num-of-cylinders

False 205

Name: num-of-cylinders, dtype: int64

engine-size

False 205

Name: engine-size, dtype: int64

fuel-system

False 205

Name: fuel-system, dtype: int64

bore

False 201

True 4

Name: bore, dtype: int64

stroke

False 201

True 4

Name: stroke, dtype: int64

compression-ratio

False 205

Name: compression-ratio, dtype: int64

horsepower

False 203

True 2

Name: horsepower, dtype: int64

peak-rpm

False 203

True 2

Name: peak-rpm, dtype: int64

city-mpg

False 205

Name: city-mpg, dtype: int64

highway-mpg

False 205

Name: highway-mpg, dtype: int64

price

False 201

```
True          4
Name: price, dtype: int64
```

Based on the summary above, each column has 205 rows of data and seven of the columns containing missing data:

“normalized-losses”: 41 missing data

“num-of-doors”: 2 missing data

“bore”: 4 missing data

“stroke” : 4 missing data

“horsepower”: 2 missing data

“peak-rpm”: 2 missing data

“price”: 4 missing data

Deal with missing data

How to deal with missing data?

Drop data a. Drop the whole row b. Drop the whole column

Replace data a. Replace it by mean b. Replace it by frequency c. Replace it based on other functions

`df.head()` Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

Replace by mean:

“normalized-losses”: 41 missing data, replace them with mean

“stroke”: 4 missing data, replace them with mean

“bore”: 4 missing data, replace them with mean

“horsepower”: 2 missing data, replace them with mean

“peak-rpm”: 2 missing data, replace them with mean

Replace by frequency:

“num-of-doors”: 2 missing data, replace them with “four”.

Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

Drop the whole row:

“price”: 4 missing data, simply delete the whole row

Reason: price is what we want to predict. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us


```
[12]: df.head()
```

```
[12]:   symboling normalized-losses      make fuel-type aspiration num-of-doors \
0         3             NaN  alfa-romero    gas      std         two
1         3             NaN  alfa-romero    gas      std         two
2         1             NaN  alfa-romero    gas      std         two
3         2            164    audi        gas      std         four
4         2            164    audi        gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0      mpfi  3.47    2.68           9.0        111    5000    21
1      mpfi  3.47    2.68           9.0        111    5000    21
2      mpfi  2.68    3.47           9.0        154    5000    19
3      mpfi  3.19    3.40          10.0        102    5500    24
4      mpfi  3.19    3.40           8.0        115    5500    18

      highway-mpg  price
0         27  13495
1         27  16500
2         26  16500
3         30  13950
4         22  17450

[5 rows x 26 columns]
```

Calculate the mean value for the “normalized-losses” column

```
[11]: avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
      print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

Replace “NaN” with mean value in “normalized-losses” column

```
[14]: df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

```
[15]: df.head()
```

```
[15]:   symboling normalized-losses      make fuel-type aspiration num-of-doors \
0         3            122.0  alfa-romero    gas      std         two
```

1	3	122.0	alfa-romero	gas	std	two
2	1	122.0	alfa-romero	gas	std	two
3	2	164	audi	gas	std	four
4	2	164	audi	gas	std	four

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

Calculate the mean value for the “bore” column

```
[16]: avg_bore=df['bore'].astype('float').mean(axis=0)
      print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810943

Replace “NaN” with the mean value in the “bore” column

```
[17]: df["bore"].replace(np.nan, avg_bore, inplace=True)
```

```
[18]: df.head()
```

```
[18]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	3	122.0	alfa-romero	gas	std	two	
1	3	122.0	alfa-romero	gas	std	two	
2	1	122.0	alfa-romero	gas	std	two	
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

Question #1:

Based on the example above, replace NaN in “stroke” column with the mean value.

```
[19]: # Write your code below and press Shift+Enter to execute
avg_stroke=df['stroke'].astype('float').mean(axis=0)
print("Average of stroke:", avg_stroke)
df["stroke"].replace(np.nan, avg_stroke, inplace=True)
```

Average of stroke: 3.255422885572139

Click here for the solution

```
#Calculate the mean vaule for "stroke" column
avg_stroke = df["stroke"].astype("float").mean(axis = 0)
print("Average of stroke:", avg_stroke)

# replace NaN by mean value in "stroke" column
df["stroke"].replace(np.nan, avg_stroke, inplace = True)
```

Calculate the mean value for the “horsepower” column

```
[20]: avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
print("Average horsepower:", avg_horsepower)
```

Average horsepower: 104.25615763546799

Replace “NaN” with the mean value in the “horsepower” column

```
[21]: df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

Calculate the mean value for “peak-rpm” column

```
[22]: avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

Replace “NaN” with the mean value in the “peak-rpm” column

```
[23]: df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the “.value_counts()” method:

```
[24]: df['num-of-doors'].value_counts()
```

```
[24]: four      114
      two       89
      Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the “.idxmax()” method to calculate the most common type automatically:

```
[25]: df['num-of-doors'].value_counts().idxmax()
```

```
[25]: 'four'
```

The replacement procedure is very similar to what we have seen previously:

```
[26]: #replace the missing 'num-of-doors' values by the most frequent
df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, let’s drop all rows that do not have price data:

```
[27]: # simply drop whole row with NaN in "price" column
df.dropna(subset=["price"], axis=0, inplace=True)

# reset index, because we dropped two rows
df.reset_index(drop=True, inplace=True)
```

```
[28]: df.head()
```

```
[28]:   symboling  normalized-losses      make fuel-type aspiration num-of-doors \
0         3         122.0  alfa-romero    gas      std         two
1         3         122.0  alfa-romero    gas      std         two
2         1         122.0  alfa-romero    gas      std         two
3         2          164      audi      gas      std         four
4         2          164      audi      gas      std         four
```

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

Good! Now, we have a dataset with no missing values.

Correct data format

We are almost there!

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use:

.dtype() to check the data type

.astype() to change the data type

Let's list the data types for each column

[29]: `df.dtypes`

```
[29]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
```

engine-location	object
wheel-base	float64
length	float64
width	float64
height	float64
curb-weight	int64
engine-type	object
num-of-cylinders	object
engine-size	int64
fuel-system	object
bore	object
stroke	object
compression-ratio	float64
horsepower	object
peak-rpm	object
city-mpg	int64
highway-mpg	int64
price	object
dtype:	object

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

Convert data types to proper format

```
[30]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

Let us list the columns after the conversion

```
[31]: df.dtypes
```

symboling	int64
normalized-losses	int64
make	object
fuel-type	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64
length	float64

```

width          float64
height         float64
curb-weight    int64
engine-type    object
num-of-cylinders object
engine-size    int64
fuel-system    object
bore           float64
stroke         float64
compression-ratio float64
horsepower     object
peak-rpm       float64
city-mpg       int64
highway-mpg    int64
price          float64
dtype: object

```

Wonderful!

Now we have finally obtained the cleaned dataset with no missing values with all data in its proper format.

Data Standardization

Data is usually collected from different agencies in different formats. (Data standardization is also a term for a particular type of data normalization where we subtract the mean and divide by the standard deviation.)

What is standardization?

Standardization is the process of transforming data into a common format, allowing the researcher to make the meaningful comparison.

Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns “city-mpg” and “highway-mpg” are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accepts the fuel consumption with L/100km standard.

We will need to apply data transformation to transform mpg into L/100km.

The formula for unit conversion is:

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

```
[32]: df.head()
```

```

[32]:   symboling  normalized-losses      make fuel-type aspiration \
0         3         122  alfa-romero      gas      std

```

1	3	122	alfa-romero	gas	std
2	1	122	alfa-romero	gas	std
3	2	164	audi	gas	std
4	2	164	audi	gas	std

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0

[5 rows x 26 columns]

```
[33]: # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
df['city-L/100km'] = 235/df["city-mpg"]

# check your transformed data
df.head()
```

```
[33]:      symboling  normalized-losses      make fuel-type aspiration \
0          3          122  alfa-romero    gas      std
1          3          122  alfa-romero    gas      std
2          1          122  alfa-romero    gas      std
3          2          164      audi    gas      std
4          2          164      audi    gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0          two  convertible      rwd      front      88.6  ...
1          two  convertible      rwd      front      88.6  ...
2          two   hatchback      rwd      front      94.5  ...
3         four      sedan      fwd      front      99.8  ...
4         four      sedan      4wd      front      99.4  ...
```


	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000.0	21	
1	mpfi	3.47	2.68	9.0	111	5000.0	21	
2	mpfi	2.68	3.47	9.0	154	5000.0	19	
3	mpfi	3.19	3.40	10.0	102	5500.0	24	
4	mpfi	3.19	3.40	8.0	115	5500.0	18	

	highway-mpg	price	city-L/100km
0	27	13495.0	11.190476
1	27	16500.0	11.190476
2	26	16500.0	12.368421
3	30	13950.0	9.791667
4	22	17450.0	13.055556

[5 rows x 27 columns]

Question #2:

According to the example above, transform mpg to L/100km in the column of “highway-mpg” and change the name of column to “highway-L/100km”.

```
[40]: # Write your code below and press Shift+Enter to execute
df["highway-mpg"] = 235/df["highway-mpg"]
```

```
[43]: df.rename(columns={'highway-mpg': 'highway-L/100km'}, inplace=True)
```

```
[44]: df.head()
```

```
[44]:      symboling  normalized-losses      make fuel-type aspiration \
0           3           122  alfa-romero      gas      std
1           3           122  alfa-romero      gas      std
2           1           122  alfa-romero      gas      std
3           2           164      audi      gas      std
4           2           164      audi      gas      std
```

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000.0	21	
1	mpfi	3.47	2.68	9.0	111	5000.0	21	
2	mpfi	2.68	3.47	9.0	154	5000.0	19	
3	mpfi	3.19	3.40	10.0	102	5500.0	24	

```
4          mpfi  3.19    3.40                8.0        115    5500.0        18
```

```
      highway-L/100km    price  city-L/100km
0          8.703704  13495.0    11.190476
1          8.703704  16500.0    11.190476
2          9.038462  16500.0    12.368421
3          7.833333  13950.0     9.791667
4          10.681818  17450.0    13.055556
```

```
[5 rows x 27 columns]
```

[Click here for the solution](#)

```
# transform mpg to L/100km by mathematical operation (235 divided by mpg)
df["highway-mpg"] = 235/df["highway-mpg"]
```

```
# rename column name from "highway-mpg" to "highway-L/100km"
df.rename(columns={'highway-mpg': 'highway-L/100km'}, inplace=True)
```

```
# check your transformed data
df.head()
```

Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling the variable so the variable values range from 0 to 1.

Example

To demonstrate normalization, let's say we want to scale the columns "length", "width" and "height".

Target: would like to normalize those variables so their value ranges from 0 to 1

Approach: replace original value by (original value)/(maximum value)

```
[45]: # replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
```

Question #3:

According to the example above, normalize the column "height".

```
[47]: # Write your code below and press Shift+Enter to execute
df['height'] = df['height']/df['height'].max()
```

```
[48]: df[["length", "width", "height"]].head()
```

```
[48]:      length      width      height
0  0.811148  0.890278  0.816054
1  0.811148  0.890278  0.816054
2  0.822681  0.909722  0.876254
3  0.848630  0.919444  0.908027
4  0.848630  0.922222  0.908027
```

[Click here for the solution](#)

```
df['height'] = df['height']/df['height'].max()

# show the scaled columns
df[["length","width","height"]].head()
```

Here we can see we’ve normalized “length”, “width” and “height” in the range of [0,1].

Binning

Why binning?

Binning is a process of transforming continuous numerical variables into discrete categorical ‘bins’ for grouped analysis.

Example:

In our dataset, “horsepower” is a real valued variable ranging from 48 to 288 and it has 59 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three ‘bins’ to simplify analysis?

We will use the pandas method ‘cut’ to segment the ‘horsepower’ column into 3 bins.

Example of Binning Data In Pandas

Convert data to correct format:

```
[49]: df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

df[["length","width","height"]].head()Let’s plot the histogram of horsepower to see what the distribution of horsepower looks like.

```
[55]: df[["horsepower"]].head()
```

```
[55]:      horsepower
0          111
1          111
2          154
3          102
4          115
```

```
[51]: %matplotlib inline
import matplotlib as plt
```


```

from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")

```

[51]: Text(0.5, 1.0, 'horsepower bins')



We would like 3 bins of equal size bandwidth so we use numpy's `linspace(start_value, end_value, numbers_generated)` function.

Since we want to include the minimum value of horsepower, we want to set `start_value = min(df["horsepower"])`.

Since we want to include the maximum value of horsepower, we want to set `end_value = max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated = 4`.

We build a bin array with a minimum value to a maximum value by using the bandwidth calculated above. The values will determine when one bin ends and another begins.

```
[53]: bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
      bins
```

```
[53]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
[57]: group_names = ['Low', 'Medium', 'High']
```

We apply the function “cut” to determine what each value of `df['horsepower']` belongs to.

```
[58]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,
      ↪include_lowest=True)
      df[['horsepower', 'horsepower-binned']].head(20)
```

```
[58]:
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low

```
[ ]:
```

Let's see the number of vehicles in each bin:

```
[59]: df["horsepower-binned"].value_counts()
```


```
[59]: Low      153
      Medium   43
      High      5
      Name: horsepower-binned, dtype: int64
```

Let's plot the distribution of each bin:

```
[60]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[60]: Text(0.5, 1.0, 'horsepower bins')
```



Look at the dataframe above carefully. You will find that the last column provides the bins for “horsepower” based on 3 categories (“Low”, “Medium” and “High”).

We successfully narrowed down the intervals from 59 to 3!

Bins Visualization

Normally, a histogram is used to visualize the distribution of bins we created above.

```
[ ]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
```

```
# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

The plot above shows the binning result for the attribute “horsepower”.

```
[61]: df.head()
```

```
[61]:   symboling  normalized-losses      make fuel-type aspiration \
0         3          122  alfa-romero      gas      std
1         3          122  alfa-romero      gas      std
2         1          122  alfa-romero      gas      std
3         2          164      audi      gas      std
4         2          164      audi      gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible      rwd      front      88.6  ...
1         two  convertible      rwd      front      88.6  ...
2         two   hatchback      rwd      front      94.5  ...
3         four      sedan      fwd      front      99.8  ...
4         four      sedan      4wd      front      99.4  ...

   bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0  3.47   2.68           9.0         111   5000.0      21
1  3.47   2.68           9.0         111   5000.0      21
2  2.68   3.47           9.0         154   5000.0      19
3  3.19   3.40          10.0         102   5500.0      24
4  3.19   3.40           8.0         115   5500.0      18

   highway-L/100km  price  city-L/100km  horsepower-binned
0         8.703704  13495.0    11.190476                Low
1         8.703704  16500.0    11.190476                Low
2         9.038462  16500.0    12.368421             Medium
3         7.833333  13950.0     9.791667                Low
4        10.681818  17450.0    13.055556                Low
```

```
[5 rows x 28 columns]
```

Indicator Variable (or Dummy Variable)

What is an indicator variable?

An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called ‘dummies’ because the numbers themselves don’t have inherent meaning.

Why we use indicator variables?

We use indicator variables so we can use categorical variables for regression analysis in the later modules.

Example

We see the column “fuel-type” has two unique values: “gas” or “diesel”. Regression doesn’t understand words, only numbers. To use this attribute in regression analysis, we convert “fuel-type” to indicator variables.

We will use pandas’ method ‘get_dummies’ to assign numerical values to different categories of fuel type.

```
[62]: df.columns
```

```
[62]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
          'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
          'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
          'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
          'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
          'highway-L/100km', 'price', 'city-L/100km', 'horsepower-binned'],
          dtype='object')
```

Get the indicator variables and assign it to data frame “dummy_variable_1”:

```
[68]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
      dummy_variable_1.head()
```

```
[68]:   diesel  gas
0         0    1
1         0    1
2         0    1
3         0    1
4         0    1
```

```
[67]: dummy_variable_1.tail()
```

```
[67]:   diesel  gas
196         0    1
197         0    1
198         0    1
199         1    0
200         0    1
```

Change the column names for clarity:

```
[90]: dummy_variable_1.rename(columns={'gas': 'fuel-type-gas', 'diesel':
    → 'fuel-type-diesel'}, inplace=True)
      dummy_variable_1.head()
```



```
[90]:      fuel-type-diesel  fuel-type-gas
0                0          1
1                0          1
2                0          1
3                0          1
4                0          1
```

In the dataframe, column 'fuel-type' has values for 'gas' and 'diesel' as 0s and 1s now.

```
[91]: df.columns
```

```
[91]: Index(['symboling', 'normalized-losses', 'make', 'aspiration', 'num-of-doors',
        'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',
        'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders',
        'engine-size', 'fuel-system', 'bore', 'stroke', 'compression-ratio',
        'horsepower', 'peak-rpm', 'city-mpg', 'highway-L/100km', 'price',
        'city-L/100km', 'horsepower-binned', 'fuel-type-diesel',
        'fuel-type-gas', 'fuel-type-diesel', 'fuel-type-gas',
        'fuel-type-diesel', 'fuel-type-gas', 'fuel-type-diesel',
        'fuel-type-gas', 'fuel-type-diesel', 'fuel-type-gas',
        'fuel-type-diesel', 'fuel-type-gas', 'fuel-type-diesel',
        'fuel-type-gas'],
        dtype='object')
```

```
[93]: df.head()
```

```
[93]:      symboling  normalized-losses      make aspiration num-of-doors \
0           3           122  alfa-romero      std          two
1           3           122  alfa-romero      std          two
2           1           122  alfa-romero      std          two
3           2           164      audi      std          four
4           2           164      audi      std          four

      body-style drive-wheels engine-location  wheel-base  length  ... \
0  convertible      rwd      front      88.6  0.811148  ...
1  convertible      rwd      front      88.6  0.811148  ...
2   hatchback      rwd      front      94.5  0.822681  ...
3      sedan      fwd      front      99.8  0.848630  ...
4      sedan      4wd      front      99.4  0.848630  ...

      fuel-type-diesel  fuel-type-gas  fuel-type-diesel  fuel-type-gas  \
0                0          1                0          1
1                0          1                0          1
2                0          1                0          1
3                0          1                0          1
4                0          1                0          1
```

	fuel-type-diesel	fuel-type-gas	fuel-type-diesel	fuel-type-gas	\
0	0	1	0	1	
1	0	1	0	1	
2	0	1	0	1	
3	0	1	0	1	
4	0	1	0	1	

	fuel-type-diesel	fuel-type-gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

[5 rows x 41 columns]

```
[88]: df[['fuel-system']].head()
```

```
[88]: fuel-system
0      mpfi
1      mpfi
2      mpfi
3      mpfi
4      mpfi
```

```
[89]: df['fuel-system'].head()
```

```
[89]: 0      mpfi
1      mpfi
2      mpfi
3      mpfi
4      mpfi
Name: fuel-system, dtype: object
```

```
[80]: # merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)
```

```
[79]: # merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)

# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis=1, inplace=True)
```

↳ -----

```

KeyError                                Traceback (most recent call
↳ last)

/tmp/ipykernel_65/3205621691.py in <module>
      3
      4 # drop original column "fuel-type" from "df"
----> 5 df.drop("fuel-type", axis=1, inplace=True)

~/conda/envs/python/lib/python3.7/site-packages/pandas/util/_decorators.
↳ py in wrapper(*args, **kwargs)
      309             stacklevel=stacklevel,
      310         )
--> 311         return func(*args, **kwargs)
      312
      313     return wrapper

~/conda/envs/python/lib/python3.7/site-packages/pandas/core/frame.py in
↳ drop(self, labels, axis, index, columns, level, inplace, errors)
      4911         level=level,
      4912         inplace=inplace,
-> 4913         errors=errors,
      4914     )
      4915

~/conda/envs/python/lib/python3.7/site-packages/pandas/core/generic.py
↳ in drop(self, labels, axis, index, columns, level, inplace, errors)
      4148         for axis, labels in axes.items():
      4149             if labels is not None:
-> 4150                 obj = obj._drop_axis(labels, axis, level=level,
↳ errors=errors)
      4151
      4152         if inplace:

~/conda/envs/python/lib/python3.7/site-packages/pandas/core/generic.py
↳ in _drop_axis(self, labels, axis, level, errors)
      4212             labels_missing = (axis.get_indexer_for(labels) ==
↳ -1).any()
      4213             if errors == "raise" and labels_missing:
-> 4214                 raise KeyError(f"{labels} not found in axis")
      4215
      4216             slicer = [slice(None)] * self.ndim

```

KeyError: "['fuel-type'] not found in axis"

```
[81]: df.head()
```

```
[81]:  symboling  normalized-losses      make aspiration num-of-doors  \
0         3             122  alfa-romero      std          two
1         3             122  alfa-romero      std          two
2         1             122  alfa-romero      std          two
3         2             164      audi      std          four
4         2             164      audi      std          four

      body-style drive-wheels engine-location  wheel-base  length  ...  \
0  convertible      rwd      front      88.6  0.811148  ...
1  convertible      rwd      front      88.6  0.811148  ...
2   hatchback      rwd      front      94.5  0.822681  ...
3      sedan      fwd      front      99.8  0.848630  ...
4      sedan      4wd      front      99.4  0.848630  ...

      fuel-type-diesel  fuel-type-gas  fuel-type-diesel  fuel-type-gas  \
0                   0              1              0              1
1                   0              1              0              1
2                   0              1              0              1
3                   0              1              0              1
4                   0              1              0              1

      fuel-type-diesel  fuel-type-gas  fuel-type-diesel  fuel-type-gas  \
0                   0              1              0              1
1                   0              1              0              1
2                   0              1              0              1
3                   0              1              0              1
4                   0              1              0              1

      fuel-type-diesel  fuel-type-gas
0                   0              1
1                   0              1
2                   0              1
3                   0              1
4                   0              1
```

[5 rows x 41 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. They're all 0s and 1s now.

Question #4:

Similar to before, create an indicator variable for the column "aspiration"

```
[94]: # Write your code below and press Shift+Enter to execute
# get indicator variables of aspiration and assign it to data frame
↳ "dummy_variable_2"
dummy_variable_2 = pd.get_dummies(df['aspiration'])
```

```
[96]: dummy_variable_2.head()
```

```
[96]:   std  turbo
0    1     0
1    1     0
2    1     0
3    1     0
4    1     0
```

```
[99]: # change column names for clarity
dummy_variable_2.rename(columns={'std':'aspiration-std', 'turbo':
↳ 'aspiration-turbo'}, inplace=True)
```

```
[100]: df.head()
```

```
[100]:   symboling  normalized-losses      make aspiration num-of-doors \
0         3          122  alfa-romero      std          two
1         3          122  alfa-romero      std          two
2         1          122  alfa-romero      std          two
3         2          164      audi      std          four
4         2          164      audi      std          four

   body-style drive-wheels engine-location  wheel-base  length  ... \
0  convertible      rwd      front      88.6  0.811148  ...
1  convertible      rwd      front      88.6  0.811148  ...
2   hatchback      rwd      front      94.5  0.822681  ...
3      sedan      fwd      front      99.8  0.848630  ...
4      sedan      4wd      front      99.4  0.848630  ...

   fuel-type-diesel  fuel-type-gas  fuel-type-diesel  fuel-type-gas  \
0                0                1                0                1
1                0                1                0                1
2                0                1                0                1
3                0                1                0                1
4                0                1                0                1

   fuel-type-diesel  fuel-type-gas  fuel-type-diesel  fuel-type-gas  \
0                0                1                0                1
1                0                1                0                1
2                0                1                0                1
3                0                1                0                1
4                0                1                0                1
```

	fuel-type-diesel	fuel-type-gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

[5 rows x 41 columns]

```
[101]: # show first 5 instances of data frame "dummy_variable_1"
dummy_variable_2.head()
```

```
[101]: aspiration-std aspiration-turbo
0          1          0
1          1          0
2          1          0
3          1          0
4          1          0
```

[Click here for the solution](#)

```
# get indicator variables of aspiration and assign it to data frame "dummy_variable_2"
dummy_variable_2 = pd.get_dummies(df['aspiration'])
```

```
# change column names for clarity
```

```
dummy_variable_2.rename(columns={'std':'aspiration-std', 'turbo': 'aspiration-turbo'}, inplace=True)
```

```
# show first 5 instances of data frame "dummy_variable_1"
dummy_variable_2.head()
```

Question #5:

Merge the new dataframe to the original dataframe, then drop the column 'aspiration'.

```
[105]: # Write your code below and press Shift+Enter to execute
# merge the new dataframe to the original dataframe
df = pd.concat([df, dummy_variable_2], axis=1)
```

```
[106]: dummy_variable_2.head()
```

```
[106]: aspiration-std aspiration-turbo
0          1          0
1          1          0
2          1          0
3          1          0
4          1          0
```

```
[107]: # drop original column "aspiration" from "df"
df.drop('aspiration', axis = 1, inplace=True)
```

```
[108]: df.head()
```

```
[108]:
```

	symboling	normalized-losses	make	num-of-doors	body-style	\
0	3	122	alfa-romero	two	convertible	
1	3	122	alfa-romero	two	convertible	
2	1	122	alfa-romero	two	hatchback	
3	2	164	audi	four	sedan	
4	2	164	audi	four	sedan	

	drive-wheels	engine-location	wheel-base	length	width	...	\
0	rwd	front	88.6	0.811148	0.890278	...	
1	rwd	front	88.6	0.811148	0.890278	...	
2	rwd	front	94.5	0.822681	0.909722	...	
3	fwd	front	99.8	0.848630	0.919444	...	
4	4wd	front	99.4	0.848630	0.922222	...	

	fuel-type-diesel	fuel-type-gas	fuel-type-diesel	fuel-type-gas	\
0	0	1	0	1	
1	0	1	0	1	
2	0	1	0	1	
3	0	1	0	1	
4	0	1	0	1	

	fuel-type-diesel	fuel-type-gas	aspiration-std	aspiration-turbo	\
0	0	1	1	0	
1	0	1	1	0	
2	0	1	1	0	
3	0	1	1	0	
4	0	1	1	0	

	aspiration-std	aspiration-turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

[5 rows x 44 columns]

[Click here for the solution](#)

```
# merge the new dataframe to the original dataframe
df = pd.concat([df, dummy_variable_2], axis=1)
```

```
# drop original column "aspiration" from "df"
```

```
df.drop('aspiration', axis = 1, inplace=True)
```

Save the new csv:

```
[109]: df.to_csv('clean_df.csv')
```

1.1.1 Thank you for completing this lab!

1.2 Author

Joseph Santarcangelo

1.2.1 Other Contributors

Mahdi Noorian PhD

Bahare Talayian

Eric Xiao

Steven Dong

Parizad

Hima Vasudevan

Fiorella Wenver

Yi Yao.

1.3 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-10-30	2.2	Lakshmi	Changed URL of csv
2020-09-09	2.1	Lakshmi	Updated Indicator Variables section
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

##

© IBM Corporation 2020. All rights reserved.